

Analysis and Implementation of Asynchronous Finite Difference Scheme for Advection - Diffusion Equation

Kumar Saurabh ¹ Prof. S. Sundar¹ Prof. Dr. Martin Frank ²

¹ Department of Mathematics
IIT Madras

² Math CCES
RWTH Aachen

August 31, 2016

Acknowledgement: DAAD, RWTH IT Center

- Many natural and engineering systems can be described with PDEs
 - Fluid mechanics, Electromagnetism, Quantum Mechanics
- Analytical Solution not known. Need to solve these problems numerically.

¹Maitham Alhubail, Qiqi Wan. The swept rule for breaking the latency barrier in time advancing PDEs

²Jagannathan & Donzis (XSEDE 2012), Sankaran et al. (SC 2012), Lee et al. (SC 2013)


- Many natural and engineering systems can be described with PDEs
 - Fluid mechanics, Electromagnetism, Quantum Mechanics
- Analytical Solution not known. Need to solve these problems numerically.
- Large number of numerical methods: Finite Difference, Finite Volume, Spectral Method, etc.
 - The problem is decomposed into a large number of Processing Element (PEs).
 - Extreme-scale computer clusters can solve PDEs using over 1,000,000 cores. ¹
 - The communication is required between the PEs to solve the PDEs to compute spatial derivatives.

¹Maitham Alhubail, Qiqi Wan. The swept rule for breaking the latency barrier in time advancing PDEs

²Jagannathan & Donzis (XSEDE 2012), Sankaran et al. (SC 2012), Lee et al. (SC 2013)

Motivation

- Many natural and engineering systems can be described with PDEs
 - Fluid mechanics, Electromagnetism, Quantum Mechanics
- Analytical Solution not known. Need to solve these problems numerically.
- Large number of numerical methods: Finite Difference, Finite Volume, Spectral Method, etc.
 - The problem is decomposed into a large number of Processing Element (PEs).
 - Extreme-scale computer clusters can solve PDEs using over 1,000,000 cores.¹
 - The communication is required between the PEs to solve the PDEs to compute spatial derivatives.
- Computation rates are much faster than communication
 - Exascale: Communication likely to be bottleneck.
 - **Direct Numerical Simulation (DNS): Communication or synchronization takes upto 50 - 70 % of the total computation time.**²
- Petascale computation: **What if one of the node fails?**



¹Maitham Alhubail, Qiqi Wan. The swept rule for breaking the latency barrier in time advancing PDEs

²Jagannathan & Donzis (XSEDE 2012), Sankaran et al. (SC 2012), Lee et al. (SC 2013)

Asynchronous Scheme for Advection Diffusion Equation

$$\frac{1}{\Delta t}(u_i^{n+1} - u_i^n) + \frac{c}{2\Delta x}(u_{i+1}^{\tilde{n}} - u_{i-1}^n) = \frac{\alpha}{\Delta x^2}(u_{i+1}^{\tilde{n}} - 2u_i^n + u_{i-1}^n) \quad (1)$$

- Regarding \tilde{n}
 - Synchronous when $\tilde{n} = n$
 - \tilde{n} can be $n, n-1, n-2, \dots$
 - Concrete value of \tilde{n} depends on hardware, network, traffic, (possible) unpredictable factors,...
 - \tilde{n} is in fact a principle random variable.

³Diego, A. Donzis and Konduri, Aditya, 2004 "Asynchronous Finite Difference Scheme for Partial Difference Equations", Journal of Computational Physics. 274(0), pp: 370-392.

⁴Mudigree, D., Sherleker, S., Ansumali, S., 2014 "Delayed Difference scheme for large scale scientific simulations", Physical Review Letters, 113(21), 218701.

⁵Assumptions and constraints stated in thesis.

Asynchronous Scheme for Advection Diffusion Equation

$$\frac{1}{\Delta t}(u_i^{n+1} - u_i^n) + \frac{c}{2\Delta x}(u_{i+1}^{\tilde{n}} - u_{i-1}^n) = \frac{\alpha}{\Delta x^2}(u_{i+1}^{\tilde{n}} - 2u_i^n + u_{i-1}^n) \quad (1)$$

- Regarding \tilde{n}
 - Synchronous when $\tilde{n} = n$
 - \tilde{n} can be $n, n-1, n-2, \dots$
 - Concrete value of \tilde{n} depends on hardware, network, traffic, (possible) unpredictable factors,...
 - \tilde{n} is in fact a principle random variable.

Analysis

- **Stability** Stable if the Synchronous Scheme is stable, irrespective of delay statistics.³
- **Truncation Error**: Not homogeneous in space and random.
 - Need for statistical description of the truncation error.
 - Strong Scaling: $\langle \bar{E} \rangle \sim O(\Delta x)$
 - Weak Scaling: $\langle \bar{E} \rangle \sim O(1)$
 - Verified by numerical experiments.
 - $\langle \bar{E} \rangle \propto \bar{k}$: Higher the delay, higher will be the error.
- **Asynchrony Tolerant Schemes**
 - Need for higher order schemes that are capable to maintain accuracy.
 - Truncation Error Analysis.
 - Previous work by Mudigree et al.⁴, Donzis and Aditya³ proposed such schemes.⁵

³Diego, A. Donzis and Konduri, Aditya, 2004 "Asynchronous Finite Difference Scheme for Partial Difference Equations", Journal of Computational Physics. 274(0), pp: 370-392.

⁴Mudigree, D., Sherleker, S., Ansumali, S., 2014 "Delayed Difference scheme for large scale scientific simulations", Physical Review Letters, 113(21), 218701.

⁵Assumptions and constraints stated in thesis.

Algorithms

- **Deterministic Asynchronous Scheme**

- Error: Deterministic - Independent of runs
- Exchange the information after a certain amount of steps. (SYNC_STEP)
- Naive way of Implementation

- **Stochastic Asynchronous Scheme**

- Error: Different for different runs.
- Do not wait for the communication to complete.
- Use the latest time values.
- Dependent on delay statistics.

Algorithms

- **Deterministic Asynchronous Scheme**
 - Error: Deterministic - Independent of runs
 - Exchange the information after a certain amount of steps. (SYNC_STEP)
 - Naive way of Implementation
- **Stochastic Asynchronous Scheme**
 - Error: Different for different runs.
 - Do not wait for the communication to complete.
 - Use the latest time values.
 - Dependent on delay statistics.

MPI Implementation

- **Two Sided Non - blocking Communication**
 - Achieved using MPI_Isend / MPI_Irecv / MPI_Test
 - Data to be transferred is stored in the buffer.
 - **Buffering** of data takes place.
 - Handshaking occurs - Matching tags and rank.
- **One Sided RMA**
 - Achieved using MPI_Put / MPI_Lock / MPI_Unlock
 - Target processor exposes its location to memory.
 - **No buffering** - Source processor writes into the target location before returning.
 - Redundancy of data is prevented by MPI_Lock.
 - No handshaking - Consent of target processor is not required.




Figure : Asynchronous Schemes




Figure : Asynchrony Tolerant (AT) Schemes

Results




Figure : Asynchronous Schemes




Figure : Asynchrony Tolerant (AT) Schemes




Figure : Effect of load and SYNC_STEP on speedup

Results




Figure : Asynchronous Schemes




Figure : Asynchrony Tolerant (AT) Schemes




Figure : Effect of load and SYNC_STEP on speedup

Deterministic Implementation with 2048 Grid Points per processor		Stochastic Implementation with 2048 Grid Points per processor	
SYNC_STEP	Speedup	Maximum Allowable Delay	Speedup
5	1.3615	5	1.8589
10	1.4474	10	2.6174
20	1.4889	20	3.7192
30	1.5300	30	6.0154

Table : Comparison of Deterministic and Stochastic Asynchronous Implementation




Figure : Asynchronous Schemes




Figure : Asynchrony Tolerant (AT) Schemes




Figure : Effect of load and SYNC_STEP on speedup

Deterministic Implementation with 2048 Grid Points per processor		Stochastic Implementation with 2048 Grid Points per processor	
SYNC_STEP	Speedup	Maximum Allowable Delay	Speedup
5	1.3615	5	1.8589
10	1.4474	10	2.6174
20	1.4889	20	3.7192
30	1.5300	30	6.0154

Table : Comparison of Deterministic and Stochastic Asynchronous Implementation

Effect of Computer Architecture on High Load

● RMA operation

Processor per node	Time (s)	Maximum Delay	Average Delay
8	204.288266	38	0.152375
4	162.511951	42	0.013750
2	122.102558	44	0.075125
1	93.667287	20	0.132000

● Two-sided Non blocking Send and Receive

Processor per node	Time (s)	Maximum Delay	Average Delay
8	151.922679	65	10.7976
4	112.257639	123	20.1246
2	111.381462	137	17.8507
1	84.915095	110	13.0063

- Original Scheme:
 - Second Order Convergence without delays.
 - First Order Convergence in presence of delays.
- Second Order Convergence in presence of delays.
 - Newer Scheme - Asynchrony Tolerant Scheme.
- Fewer Communication calls \Rightarrow More speedup.
- Higher Load on processor \Rightarrow Communication cost less significant.
- Stochastic Implementation - More Speedup.
- Study of Delay Statistics.
 - MPI_RMA: better for asynchronous case in terms of average delay statistics.
 - Dependence of load on the processor. More frequent communication \Rightarrow Larger Delays.
 - Impact of Computer Architecture.